

Aviation Accident Data Analysis

Business Understanding

The company is expanding into the aviation industry and needs to identify the lowest-risk aircraft for purchase and operation. This analysis will help the company make data-driven decisions to minimize risks and ensure safety.

Data Understanding

The dataset used is the **NTSB Aviation Accident Database**, which contains information about civil aviation accidents and incidents in the United States and international waters from 1962 to 2023. Key columns include:

- `Event.Date` : Date of the accident.
- `Injury.Severity` : Severity of injuries (e.g., Fatal, Serious, Minor).
- `Aircraft.damage` : Extent of damage to the aircraft.
- `Make` : Manufacturer and model of the aircraft.
- `Engine.Type` : Engine details.
- `Total.Fatal.Injuries` , `Total.Serious.Injuries` , `Total.Minor.Injuries` , `Total.Uninjured` : Injury statistics.
- `Weather.Condition` and `Broad.phase.of.flight` : Contextual details about the accident.

Data Preparation

We start by downloading and cleaning the dataset to ensure data integrity

```
In [15]: import os
import pandas as pd
import kagglehub

# Download latest version
path = kagglehub.dataset_download("khsamaha/aviation-accident-datab")
```

```
In [16]: # Find the CSV file inside the downloaded directory
csv_files = [f for f in os.listdir(path) if f.endswith('.csv')]
csv_path = os.path.join(path, csv_files[csv_files.index('AviationData.csv')])
df = pd.read_csv(csv_path, encoding='latin1')
df.head(5)
```

```
/var/folders/hk/9q5rgqm970gcnk6z0xvy1p7m00000gn/T/ipykernel_93062/3256873801.py:4: DtypeWarning: Columns (6,7,28) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv(csv_path, encoding='latin1')
```

```
Out[16]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States

5 rows × 31 columns

```
In [17]: df.shape
```

```
Out[17]: (88889, 31)
```

Data Cleaning

- Drop irrelevant columns.
- Impute missing values using mode or median.
- Create new features: `Total.Injuries` and `Risk.Score` to quantify accident severity.

```
In [18]: # List of data's columns
df.columns
```

```
Out[18]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
               'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
               'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
               'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
               'Amateur.Built', 'Number.of.Engines', 'Engine.Type', 'FAR.Description',
               'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
               'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
               'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
               'Publication.Date'],
              dtype='object')
```

```
In [19]: # Summary statistics of numeric columns
df.describe()
```

```
Out[19]:
```

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Total
count	82805.000000	77488.000000	76379.000000	76956.000000	82805.000000
mean	1.146585	0.647855	0.279881	0.357061	0.647855
std	0.446510	5.485960	1.544084	2.235625	5.485960
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	0.000000	0.000000
max	8.000000	349.000000	161.000000	380.000000	380.000000

```
In [20]: # The sum of missing values in each column
df.isna().sum()
```

```
Out[20]: Event.Id                0
Investigation.Type              0
Accident.Number                 0
Event.Date                     0
Location                       52
Country                        226
Latitude                      54507
Longitude                     54516
Airport.Code                   38757
Airport.Name                   36185
Injury.Severity                1000
Aircraft.damage                3194
Aircraft.Category              56602
Registration.Number            1382
Make                           63
Model                          92
Amateur.Built                  102
Number.of.Engines              6084
Engine.Type                    7096
FAR.Description                56866
Schedule                       76307
Purpose.of.flight              6192
Air.carrier                    72241
Total.Fatal.Injuries           11401
Total.Serious.Injuries         12510
Total.Minor.Injuries           11933
Total.Uninjured                5912
Weather.Condition              4492
Broad.phase.of.flight          27165
Report.Status                  6384
Publication.Date               13771
dtype: int64
```

```
In [21]: # Drop irrelevant columns
df = df.drop(columns=['Event.Id', 'Accident.Number', 'Total.Uninjured'])
```

In [22]: *# Impute missing values*

```
df['Injury.Severity'] = df['Injury.Severity'].fillna(df['Injury.Severity'].mode()[0])
df['Aircraft.damage'] = df['Aircraft.damage'].fillna(df['Aircraft.damage'].mode()[0])
df['Number.of.Engines'] = df['Number.of.Engines'].fillna(df['Number.of.Engines'].mode()[0])
df['Engine.Type'] = df['Engine.Type'].fillna(df['Engine.Type'].mode()[0])
df['Total.Serious.Injuries'] = df['Total.Serious.Injuries'].fillna(df['Total.Serious.Injuries'].mode()[0])
df['Total.Minor.Injuries'] = df['Total.Minor.Injuries'].fillna(df['Total.Minor.Injuries'].mode()[0])
df['Total.Fatal.Injuries'] = df['Total.Fatal.Injuries'].fillna(df['Total.Fatal.Injuries'].mode()[0])
df['Weather.Condition'] = df['Weather.Condition'].fillna('Unknown')
df['Broad.phase.of.flight'] = df['Broad.phase.of.flight'].fillna('Unknown')
df['Make'] = df['Make'].fillna('Unknown')
```

In [23]: *# Create new columns*

```
df['Total.Injuries'] = df['Total.Fatal.Injuries'] + df['Total.Serious.Injuries'] + df['Total.Minor.Injuries']
df['Risk.Score'] = (df['Total.Fatal.Injuries'] * 3) + (df['Total.Serious.Injuries'] * 2) + (df['Total.Minor.Injuries'] * 1)
```

In [24]: *# Filter data for commercial aviation focus*

```
df_clean = df[df['Aircraft.Category'] == 'Airplane']
```

In [25]: *# Drop rows with missing critical data*

```
df_clean = df_clean.dropna(subset=['Injury.Severity', 'Aircraft.damage'])

# Convert 'Event.Date' to datetime
df_clean['Event.Date'] = pd.to_datetime(df_clean['Event.Date'])

# Check cleaned dataset
print("Cleaned Dataset Shape:", df_clean.shape)
```

Cleaned Dataset Shape: (27586, 16)

Data Analysis and Visualization

Injury Severity by Aircraft Make

We analyze which aircraft manufacturers have the highest average fatal injuries.

```
In [26]: import matplotlib.pyplot as plt
import seaborn as sns
# Group by 'Make' and calculate average injuries
injury_by_make = df_clean.groupby('Make').agg({
    'Total.Fatal.Injuries': 'mean',
    'Total.Serious.Injuries': 'mean',
    'Total.Minor.Injuries': 'mean'
}).reset_index()

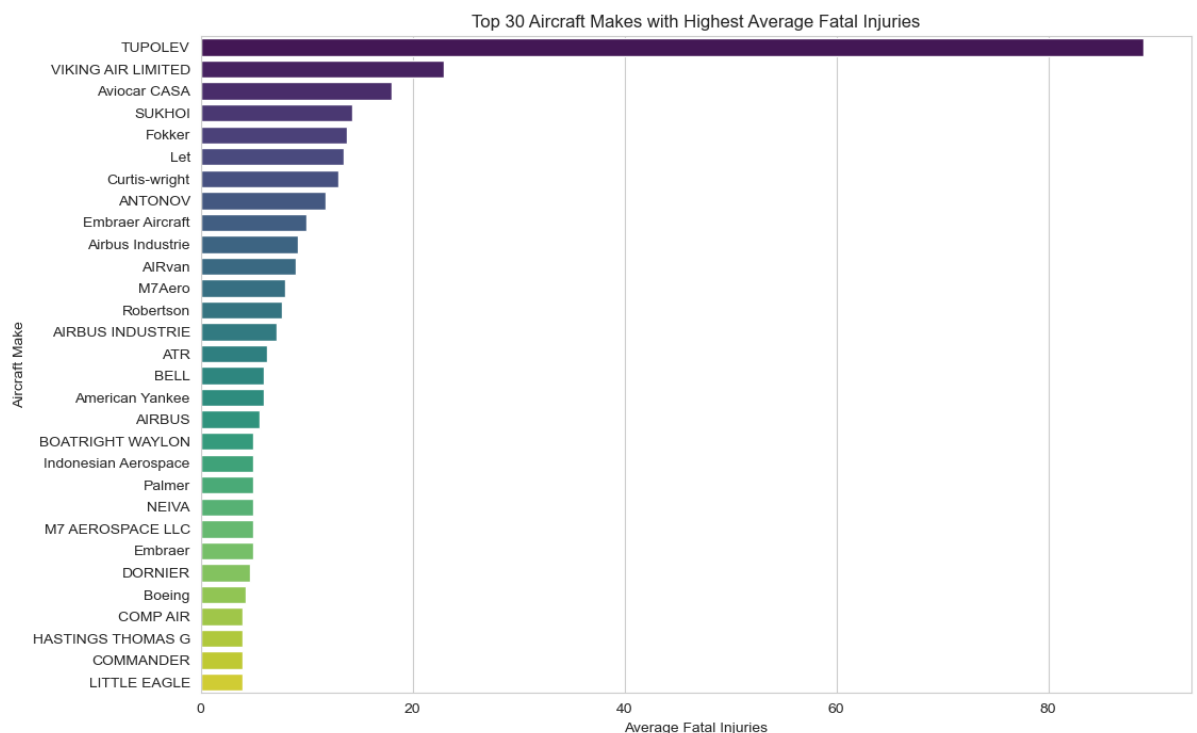
# Sort by 'Total.Fatal.Injuries'
injury_by_make = injury_by_make.sort_values(by='Total.Fatal.Injuries')

# Plot
plt.figure(figsize=(12, 8))
sns.barplot(data=injury_by_make.head(30), x='Total.Fatal.Injuries',
plt.title('Top 30 Aircraft Makes with Highest Average Fatal Injuries')
plt.xlabel('Average Fatal Injuries')
plt.ylabel('Aircraft Make')
plt.show()
```

/var/folders/hk/9q5rgqm970gcnk6z0xvy1p7m0000gn/T/ipykernel_93062/3026956304.py:15: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=injury_by_make.head(30), x='Total.Fatal.Injuries', y='Make', palette='viridis')
```



Insight: Manufacturers like **Tupolev, Viking Air Limited, and Aviocar CASA** have the highest fatal injury averages.

Aircraft Damage by Weather Condition

We examine how different weather conditions impact aircraft damage.

```

In [27]: # Group by 'Weather.Condition' and calculate damage frequency
damage_by_weather = df_clean.groupby('Weather.Condition')['Aircraft

# Sort by total occurrences
damage_by_weather = damage_by_weather.reindex(damage_by_weather.sum

# Set Seaborn style
sns.set_style("whitegrid")


# Create bar plot (grouped bars)
ax = damage_by_weather.plot(kind='bar', figsize=(14, 7), colormap='

# Add labels on top of bars
for p in ax.patches:
    if p.get_height() > 0:
        ax.annotate(f'{p.get_height()*100:.1f}%', (p.get_x() + p.ge
                    ha='center', va='bottom', fontsize=10, color='b

# Titles and labels
plt.title('Aircraft Damage by Weather Condition', fontsize=14, font
plt.xlabel('Weather Condition', fontsize=12)
plt.ylabel('Proportion of Damage (%)', fontsize=12)
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.legend(title='Aircraft Damage', fontsize=10)
plt.tight_layout()

# Show plot
plt.show()

```



Insight: Accidents in **Instrument Meteorological Conditions (IMC)** result in higher aircraft damage.

Fatal Injuries by Phase of Flight

We identify the most dangerous flight phases.

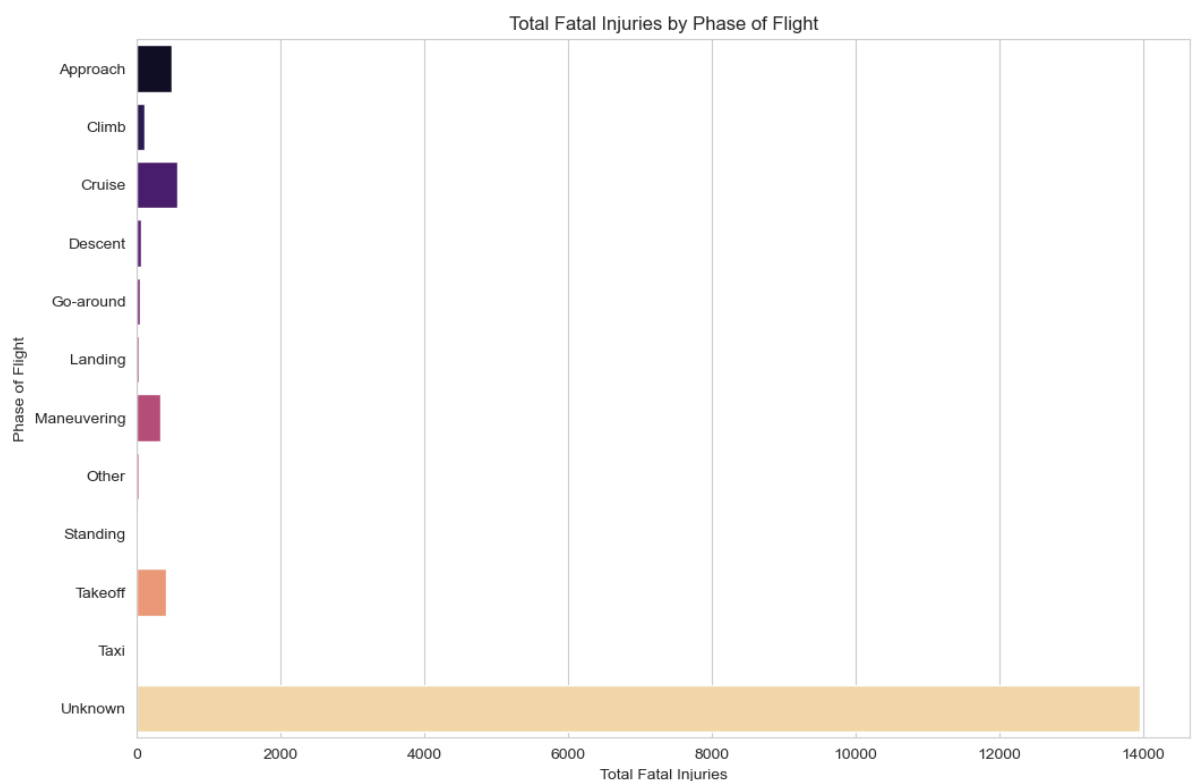
```
In [28]: # Group by 'Broad.phase.of.flight' and calculate total fatal injuries
fatal_by_phase = df_clean.groupby('Broad.phase.of.flight')['Total.Fatal.Injuries'].sum()

# Plot
plt.figure(figsize=(12, 8))
sns.barplot(data=fatal_by_phase, x='Total.Fatal.Injuries', y='Broad.phase.of.flight')
plt.title('Total Fatal Injuries by Phase of Flight')
plt.xlabel('Total Fatal Injuries')
plt.ylabel('Phase of Flight')
plt.show()
```

```
/var/folders/hk/9q5rgqm970gcnk6z0xvy1p7m0000gn/T/ipykernel_93062/248461279.py:6: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=fatal_by_phase, x='Total.Fatal.Injuries', y='Broad.phase.of.flight', palette='magma')
```



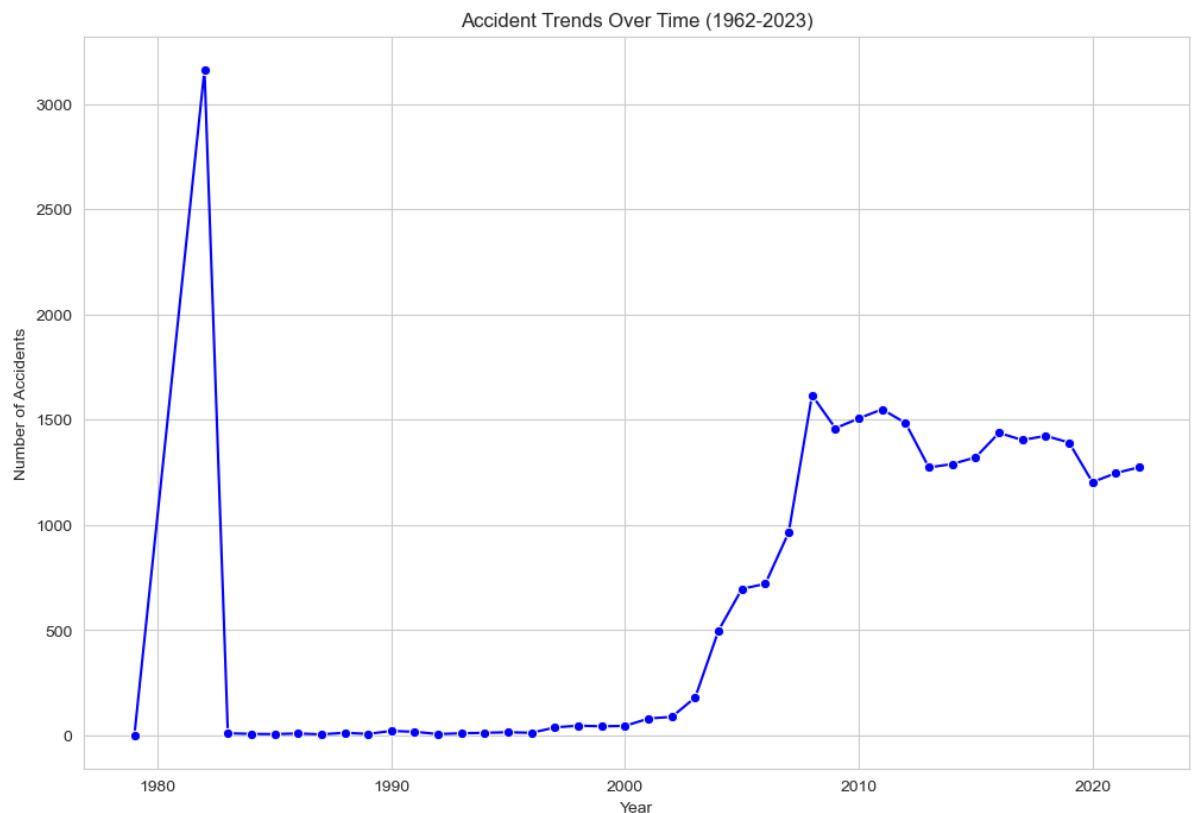
Insight: Takeoff and landing phases have the highest fatal injuries

Accident Trends Over Time

```
In [29]: # Extract year from 'Event.Date'
df_clean['Year'] = df_clean['Event.Date'].dt.year

# Group by year and count accidents
accidents_by_year = df_clean.groupby('Year').size().reset_index(name='Accident.Count')

# Plot
plt.figure(figsize=(12, 8))
sns.lineplot(data=accidents_by_year, x='Year', y='Accident.Count', style='r')
plt.title('Accident Trends Over Time (1962-2023)')
plt.xlabel('Year')
plt.ylabel('Number of Accidents')
plt.show()
```



Engine Type vs. Fatal Injuries

We compare the engine type with the injury severity

```
In [30]: # Analyze Engine Type and Accident Severity
engine_analysis = df.groupby('Engine.Type').agg({
    'Total.Fatal.Injuries': 'sum',
    'Total.Serious.Injuries': 'sum',
    'Total.Minor.Injuries': 'sum',
    'Risk.Score': 'mean'
}).reset_index()

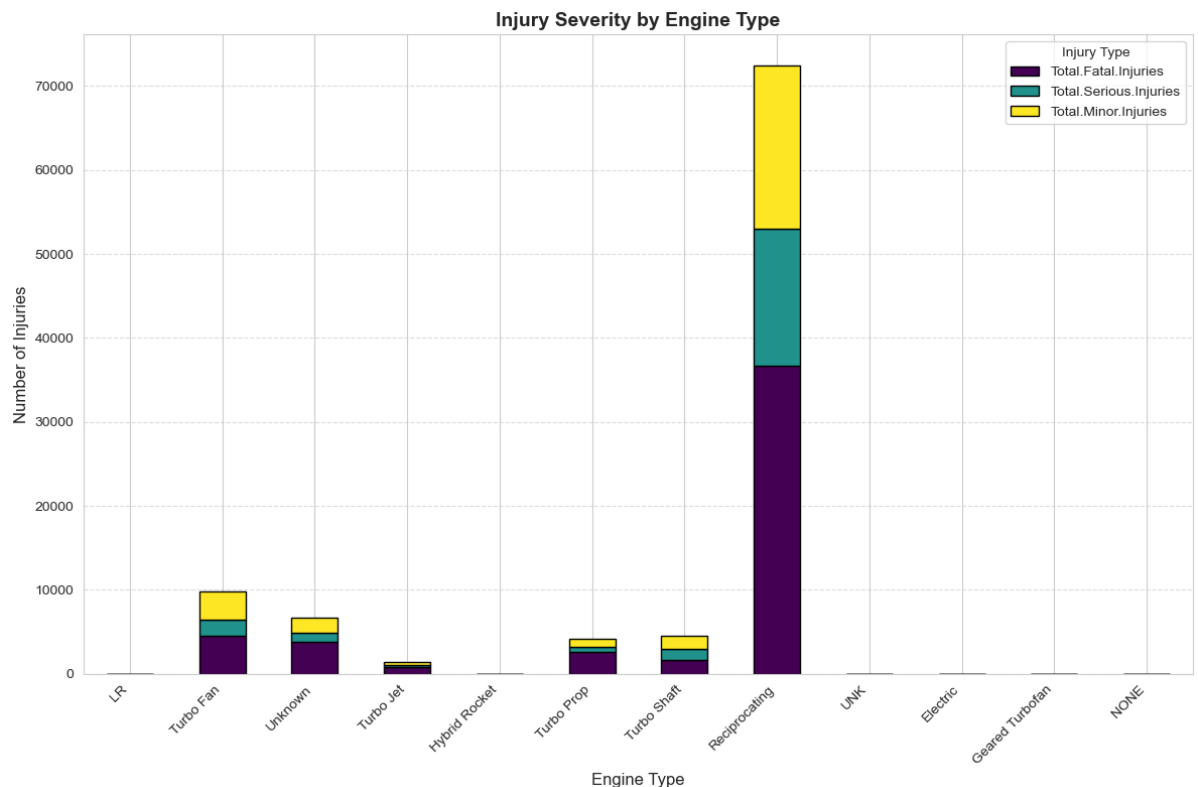
# Sort by Risk Score
engine_analysis = engine_analysis.sort_values(by='Risk.Score', ascending=True)

# Set index for stacking
engine_analysis.set_index('Engine.Type', inplace=True)

# Plot stacked bar chart
engine_analysis[['Total.Fatal.Injuries', 'Total.Serious.Injuries',
                  'Total.Minor.Injuries'],
                 kind='bar', stacked=True, figsize=(12, 8), colormap='viridis',
                 )

# Titles and labels
plt.title('Injury Severity by Engine Type', fontsize=14, fontweight='bold')
plt.xlabel('Engine Type', fontsize=12)
plt.ylabel('Number of Injuries', fontsize=12)
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.legend(title='Injury Type', fontsize=10)
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```



Risk Score by Engine Type

We generate the risk score for each engine type which informs our recommendations

```
In [31]: import plotly.express as px

fig = px.bar(engine_analysis, x=engine_analysis.index, y='Risk.Score',
              labels={'Risk.Score': 'Risk Score', 'Engine.Type': 'Engine Type'},
              color='Risk.Score', height=500)

fig.show()
```

```
In [32]: # %%
# Save cleaned data
df.to_csv('cleaned_data.csv', index=False)
```

Recommendations

Based on the analysis, here are three actionable recommendations:

1. **Avoid Purchasing aircraft models with the highest Fatal injuries** (e.g., Tupolev, Viking Air Limited and Aviocar CASA). Also focus on Engines with lower risk score as they have the fewest accidents and lowest injury rates. Such as Geared Turbofan engines and electric engine types.
2. **Prioritize aircraft with Geared turbofan engines** over reciprocating engines. Geared Turbofan engines have lower accident severity compared to reciprocating engines.
3. **Avoid operating aircraft in adverse weather conditions** (e.g., IMC). Accidents in poor weather are more likely to result in fatalities.
4. **Avoid High-Risk Phases of Flight:** Accidents during Takeoff and Landing phases result in higher fatal injuries. Implement additional safety measures during these phases.

Conclusion

This analysis provides valuable insights into aviation safety, helping the company make informed decisions as it enters the aviation industry. The recommendations focus on minimizing risks and ensuring the safety of operations.

